

# MxNet Gluon

**Basics, Computer Vision, NLP (and even more NLP)**  
**Part I (Basics)**

**Leonard Lausen  
Haibin Lin  
Alex Smola**

# Outline

<b>8:30-9:15</b>	Installation and Basics (NDArray, AutoGrad, Libraries)
<b>9:15-9:30</b>	Neural Networks 101 (MLP, ConvNet, LSTM, Loss, SGD) - Part I
<b>9:30-10:00</b>	Break
<b>10:00-10:30</b>	Neural Networks 101 (MLP, ConvNet, LSTM, Loss, SGD) - Part II
<b>10:30-11:00</b>	Computer Vision 101 (Gluon CV)
<b>11:00-11:30</b>	Parallel and distributed training
<b>11:30-12:00</b>	Data I/O in NLP (and iterators)
<b>12:00-13:30</b>	Break
<b>13:30-14:15</b>	Embeddings
<b>14:15-15:00</b>	Language models (LM)
<b>15:00-15:30</b>	Sequence Generation from LM
<b>15:30-16:00</b>	Break
<b>16:00-16:15</b>	Sentiment analysis
<b>16:15-17:00</b>	Transformer Models & machine translation
<b>17:00-17:30</b>	Questions



# What can deep learning do?

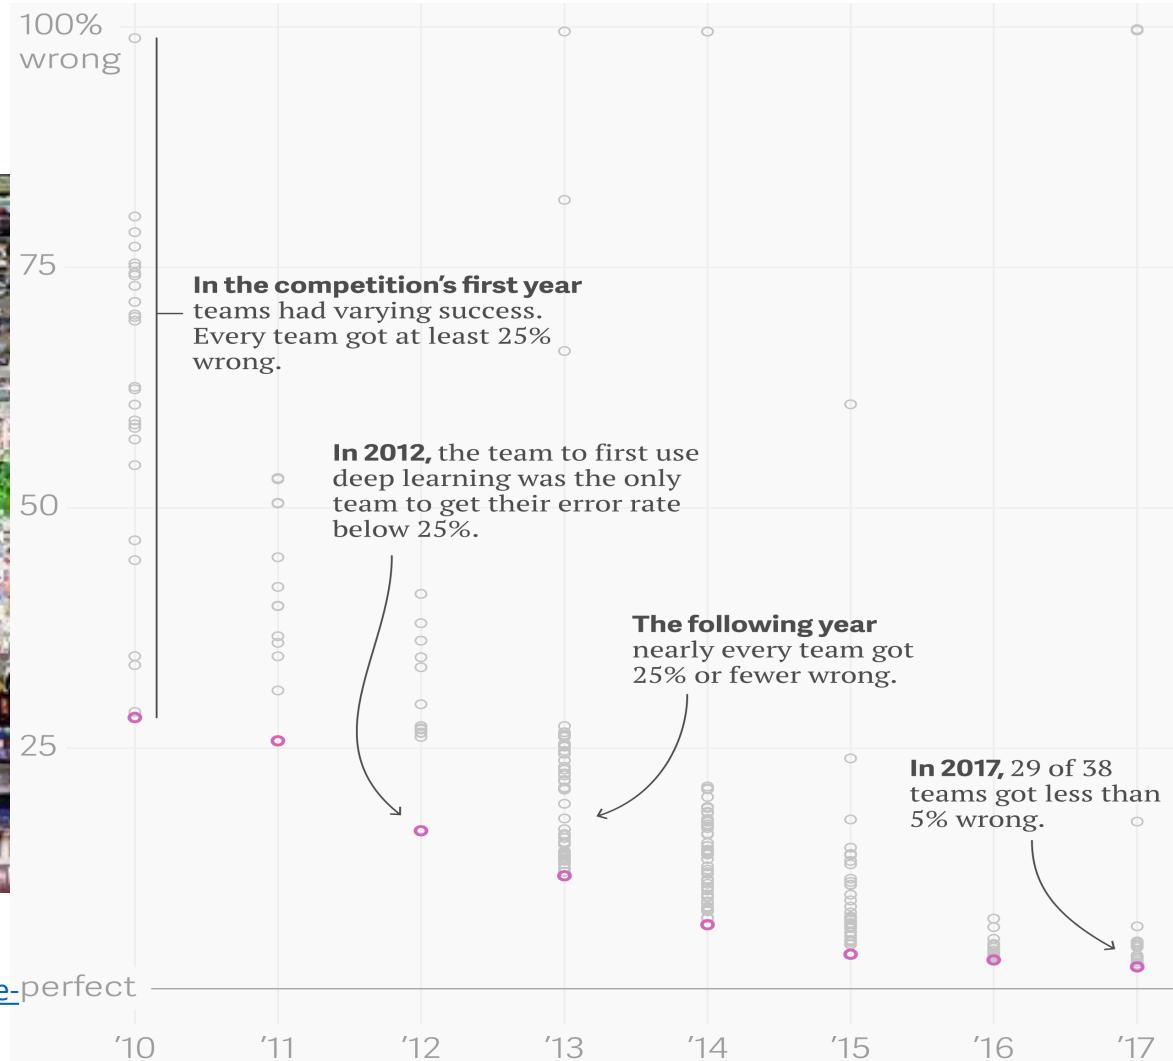
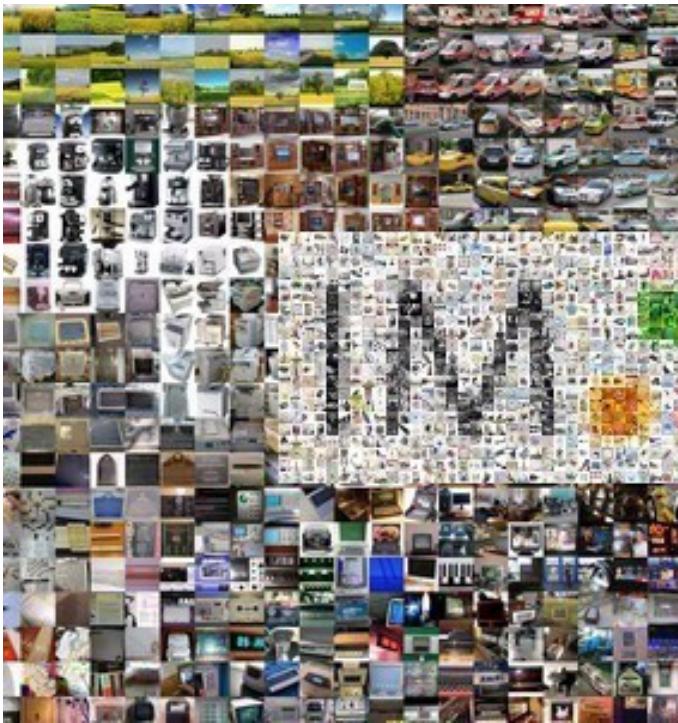
# Classify Images



<http://www.image-net.org/>



# Classify Images



Yanofsky, Quartz

<https://qz.com/1034972/the-data-that-changed-the-perfect-direction-of-ai-research-and-possibly-the-world/>

# Detect and Segment Objects



[https://github.com/matterport/Mask\\_RCNN](https://github.com/matterport/Mask_RCNN)



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



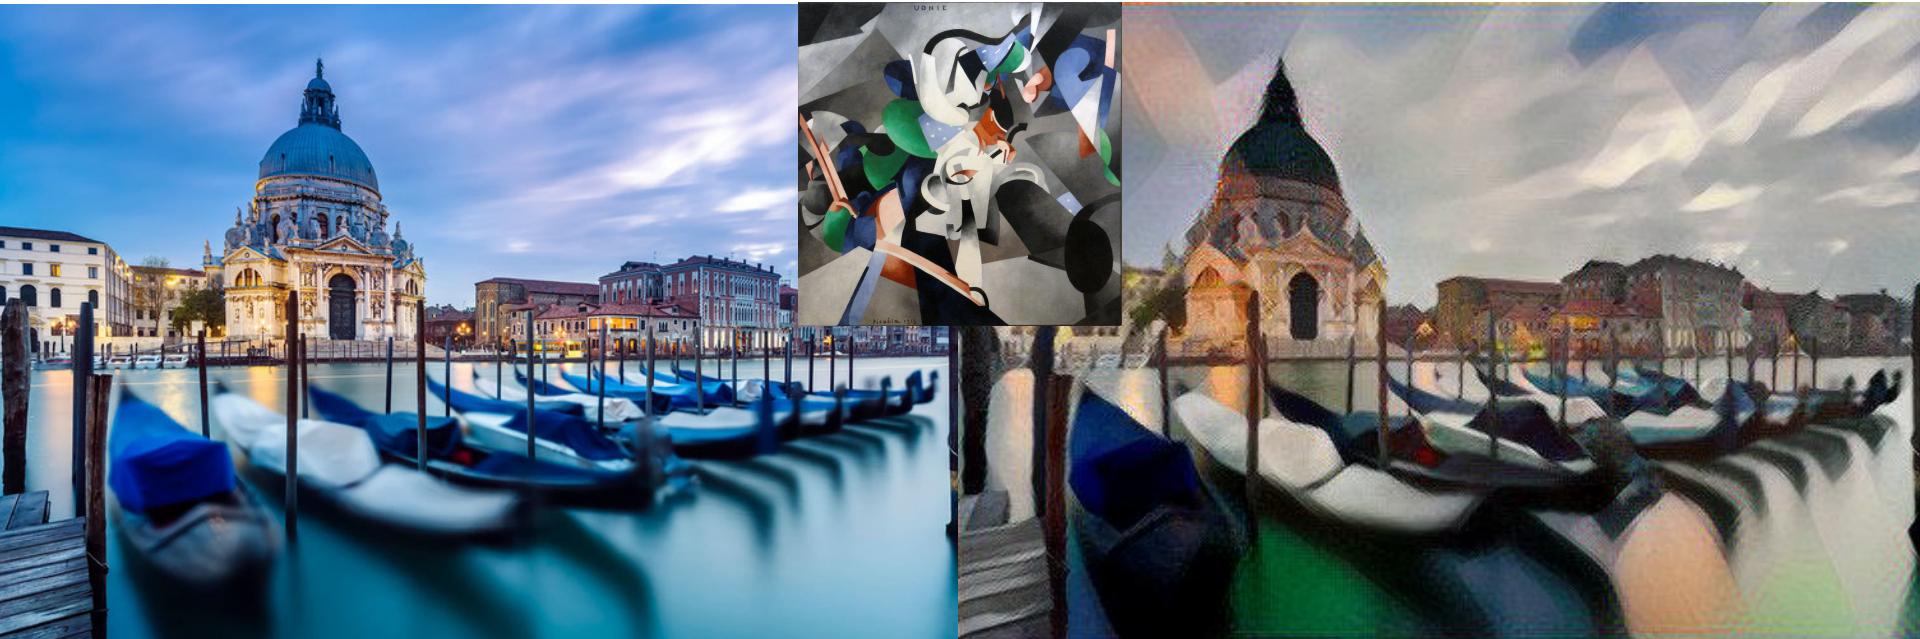
# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



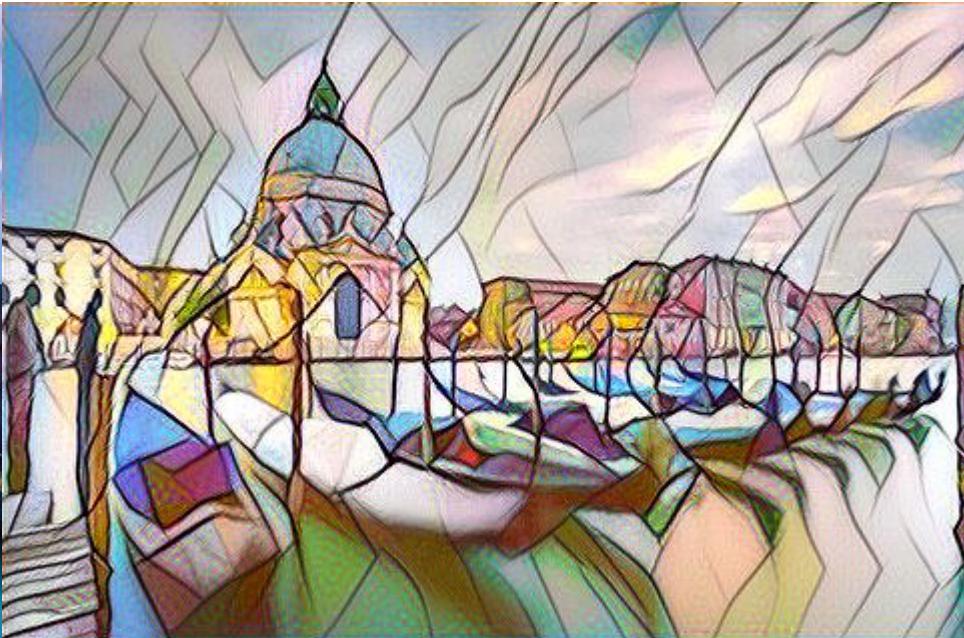
# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



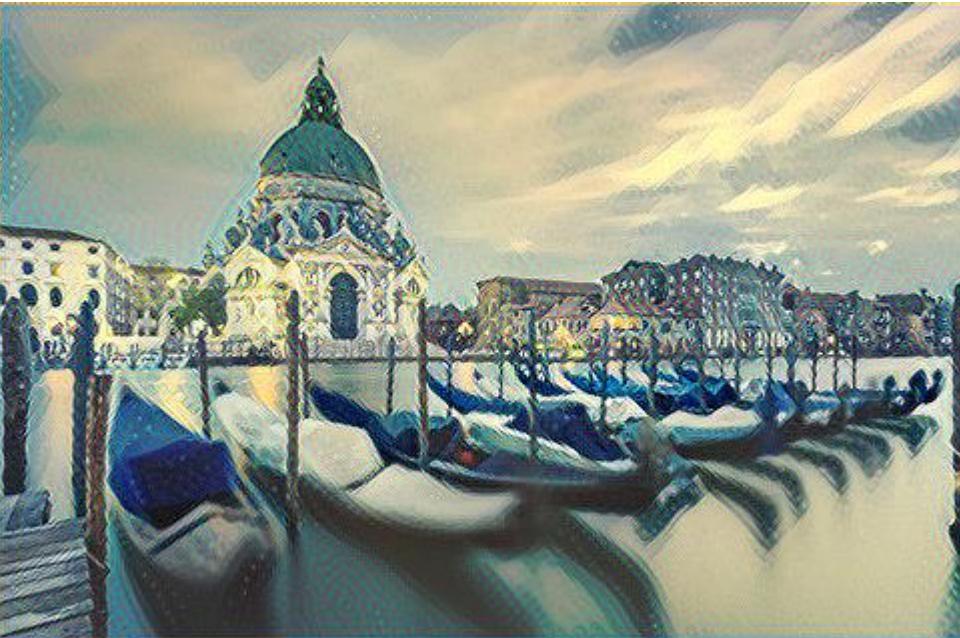
# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



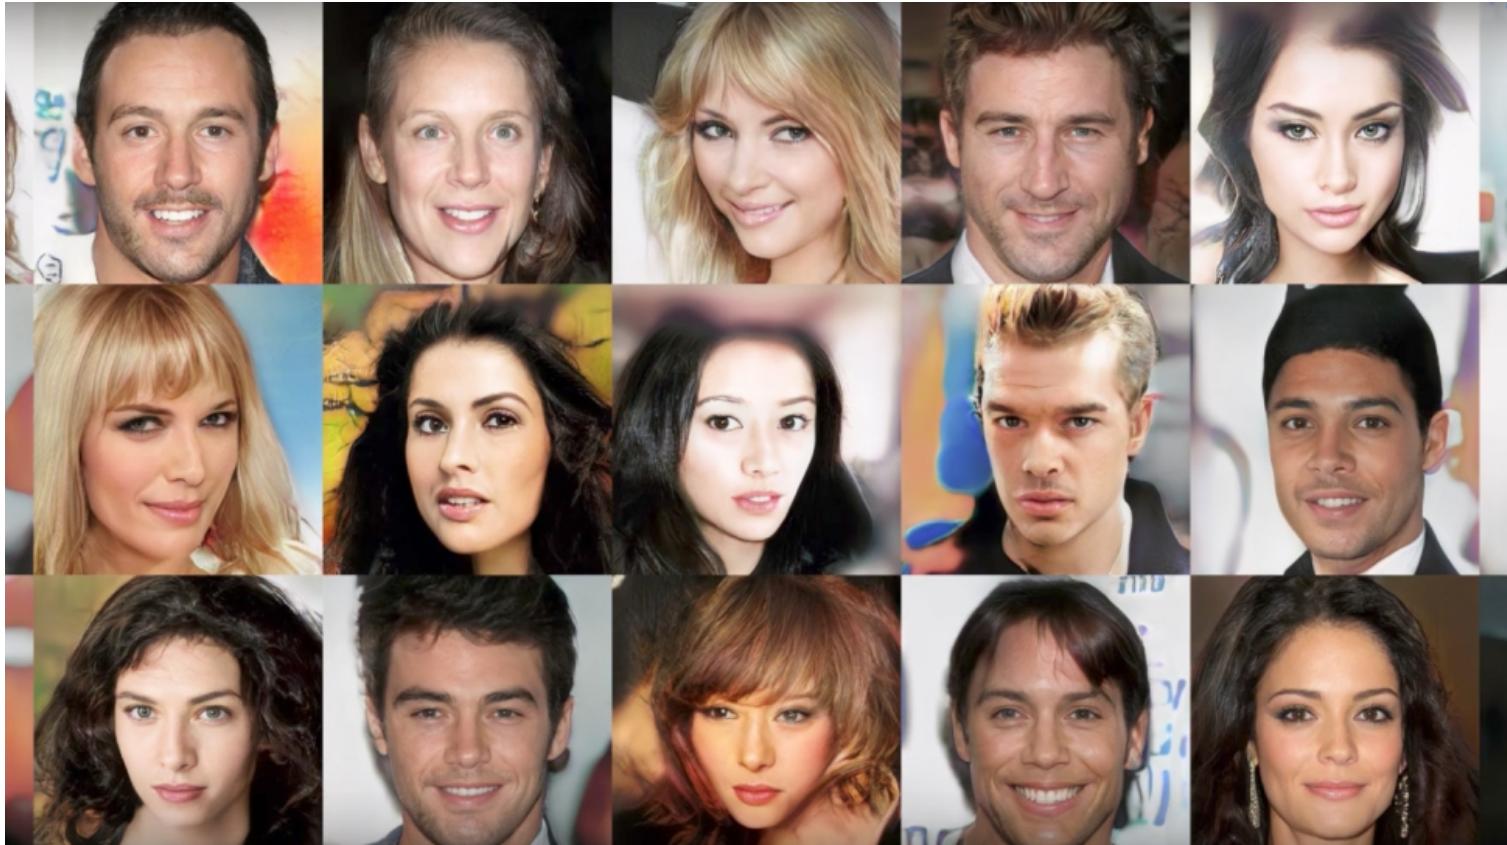
# Style transfer



<https://github.com/zhanghang1989/MXNet-Gluon-Style-Transfer/>



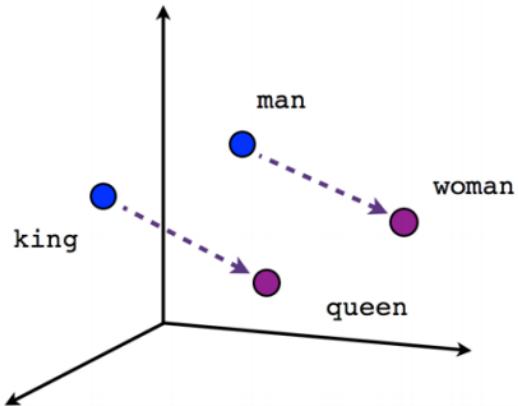
# Synthesize Faces



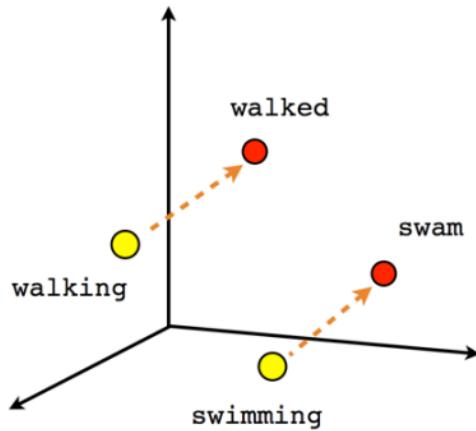
Karras et al, ICLR 2018

aws

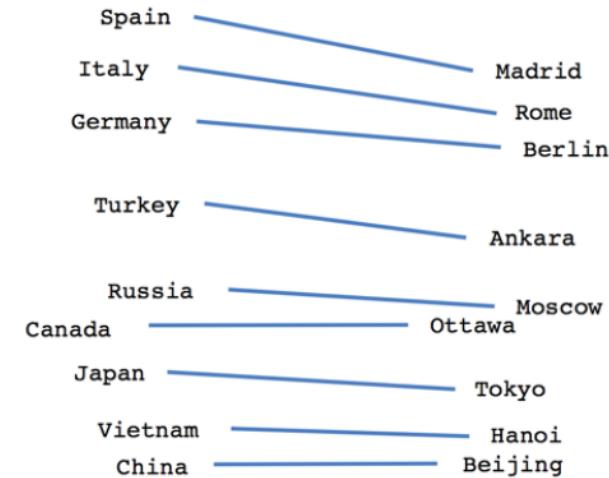
# Analogies



Male-Female



Verb tense

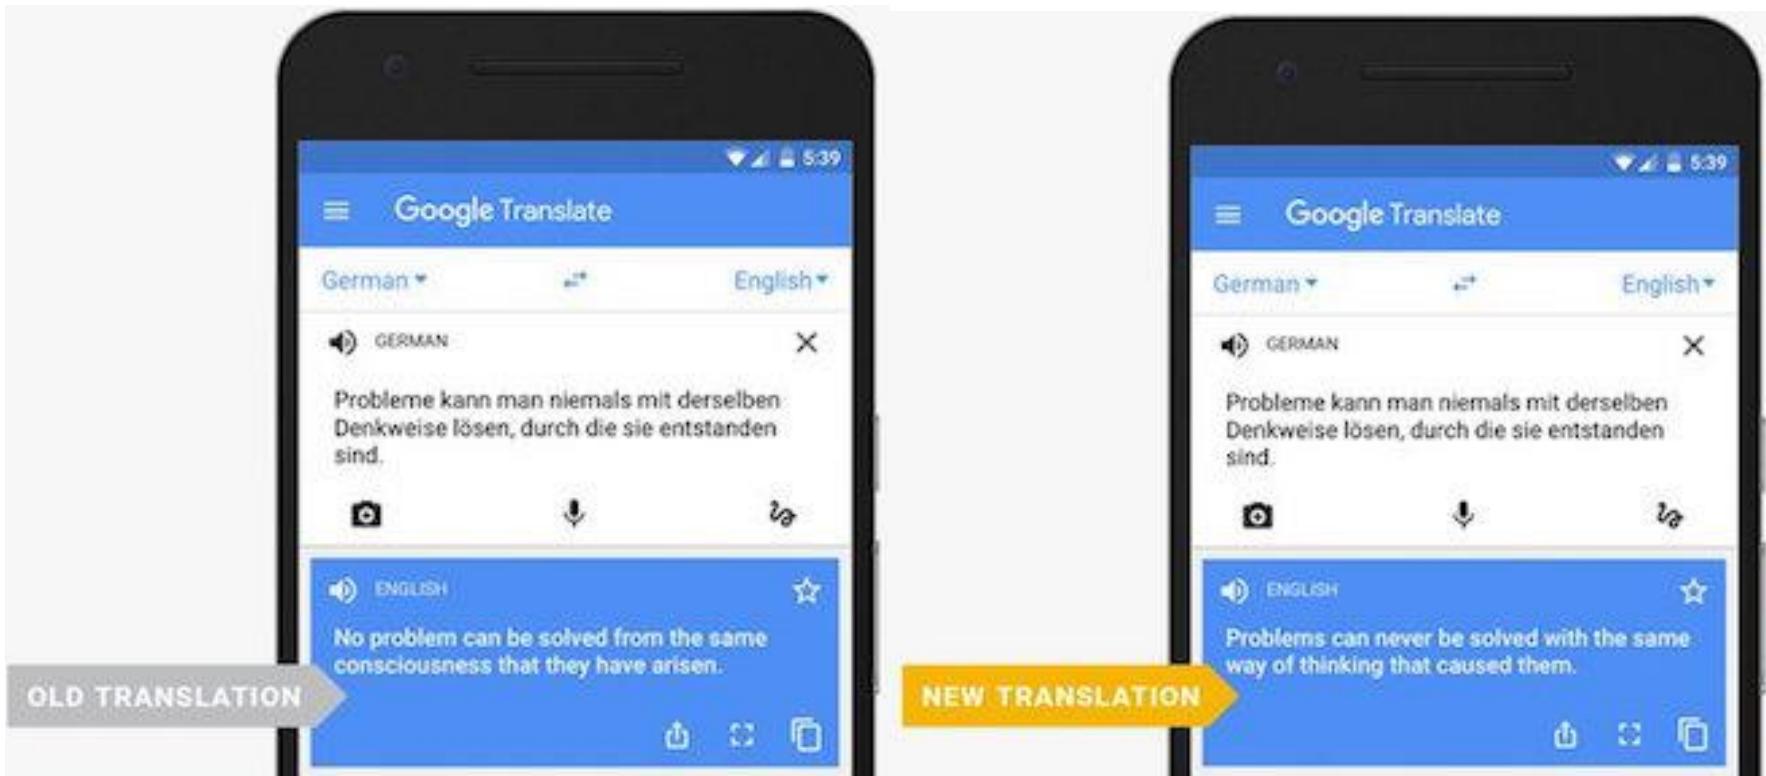


Country-Capital

<https://www.tensorflow.org/tutorials/word2vec>



# Machine Translation



<https://www.pcmag.com/news/349610/google-expands-neural-networks-for-language-translation>



# Text synthesis

**Content:** Two dogs play by a tree.

**Style:** **happily, love**



Two dogs **in love** play **happily** by a tree.

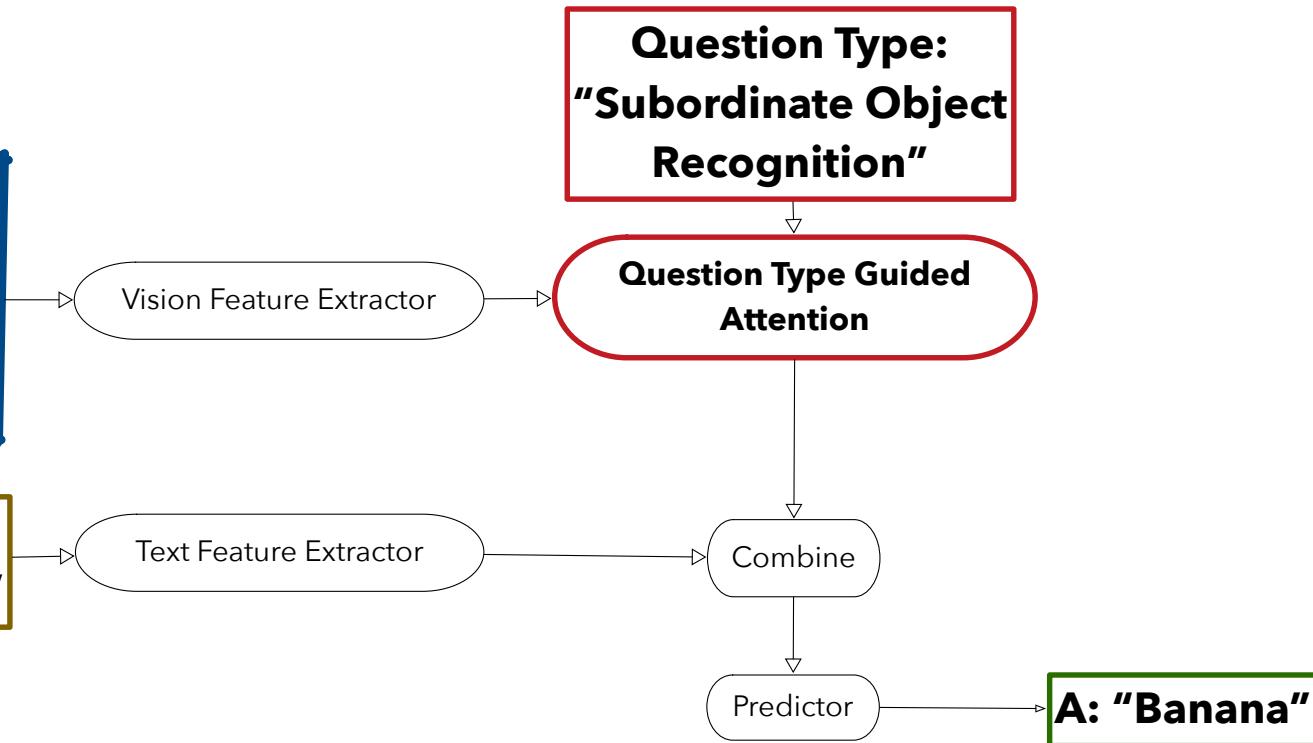
Li et al, NACCL, 2018



# Question answering



**Q: "What's her  
mustache made of?"**

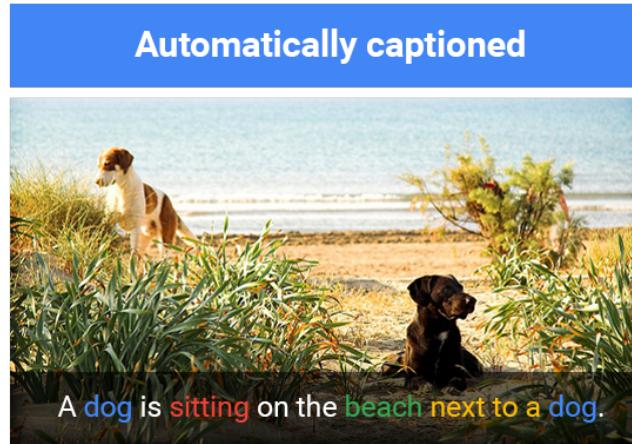
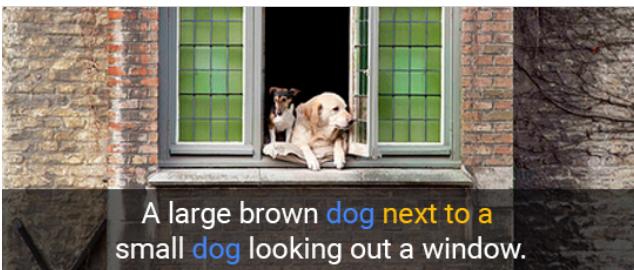


Shi et al, 2018, Arxiv



# Image captioning

## Human captions from the training set



Shallue et al, 2016

<https://ai.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>



**OK, I'm sold ... I want this! Now!**



# Programming for deep learning used to be...

```
void computeLogregSoftmaxGrad(NVMatrix& labels, NVMatrix& probs, NVMatrix& target, bool add, float coeff) {  
    int numCases = probs.getLeadingDim();  
    int numOut = probs.getFollowingDim();  
    assert(labels.getNumElements() == numCases);  
    assert(probs.isContiguous());  
    assert(target.isContiguous());  
    assert(labels.isContiguous());  
    assert(probs.isTrans());  
  
    dim3 threads(LOGREG_GRAD_THREADS_X, LOGREG_GRAD_THREADS_Y);  
    dim3 blocks(DIVUP(numCases, LOGREG_GRAD_THREADS_X), DIVUP(numOut, LOGREG_GRAD_THREADS_Y));  
    if (!add) {  
        target.resize(probs);  
        kLogregSoftmaxGrad<false><<<blocks, threads>>>(probs.getDevData(), labels.getDevData(), target.getDevData(),  
                                                               numCases, numOut, coeff);  
    } else {  
        kLogregSoftmaxGrad<true><<<blocks, threads>>>(probs.getDevData(), labels.getDevData(), target.getDevData(),  
                                                               numCases, numOut, coeff);  
    }  
  
    getLastCudaError("computeLogregSoftmaxGrad: Kernel execution failed");  
}
```



# Programming for deep learning today

```
loss_function = gluon.loss.SoftmaxCrossEntropyLoss()  
loss = loss_function(output)  
loss.backward()
```



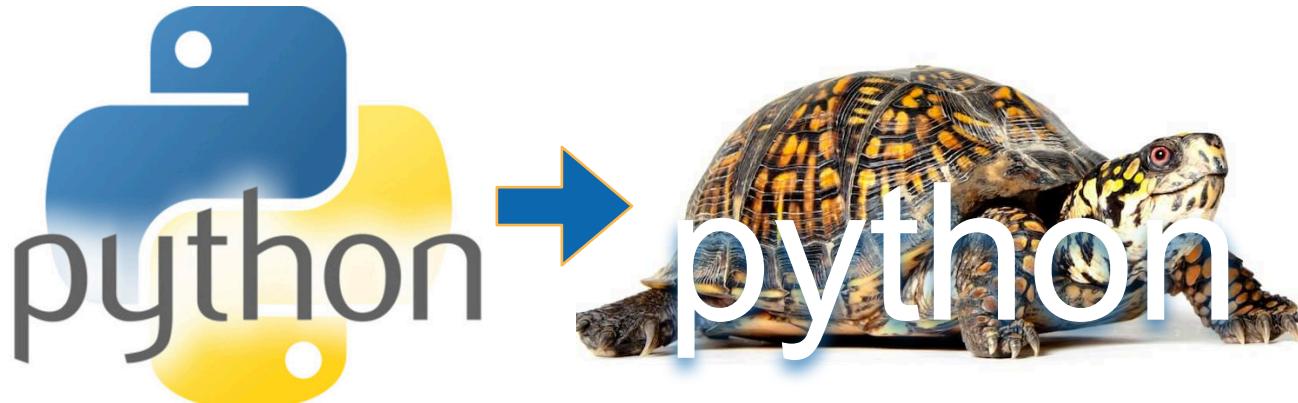
# Programming for deep learning today

```
loss_function = gluon.loss.SoftmaxCrossEntropyLoss()  
loss = loss_function(output)  
loss.backward()
```

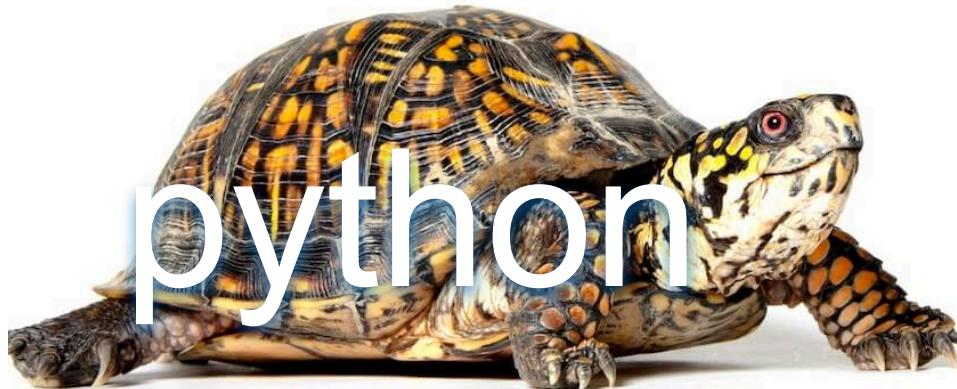


# Programming for deep learning today

```
loss_function = gluon.loss.SoftmaxCrossEntropyLoss()  
loss = loss_function(output)  
loss.backward()
```



# Programming for deep learning today



**PYTORCH**



**theano**

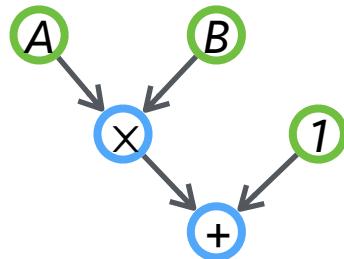
```
output = model(data)
```



[crazy GPU stuff]

**aws**

# Declarative (Symbolic) Programs



```
A = Variable('A')  
B = Variable('B')  
C = B * A  
D = C + 1  
f = compile(D)  
d = f(A=np.ones(10),  
      B=np.ones(10)*2)
```

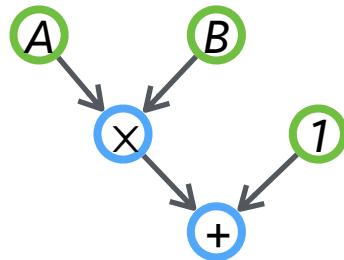
- **Advantages:**

- Opportunities for optimization
- Easy to serialize models
- More portable across languages

- **Disadvantages:**

- Hard to debug
- Unsuitable for dynamic graphs
- Can't use native code

# Declarative (Symbolic) Programs



```
A = Variable('A')  
B = Variable('B')  
C = B * A  
D = C + 1  
f = compile(D)  
d = f(A=np.ones(10),  
      B=np.ones(10)*2)
```

- **Advantages:**

- Opportunities for optimization
- Easy to serialize models
- More portable across languages

- **Disadvantages:**

- Hard to debug
- Unsuitable for dynamic graphs
- Can't use native code

C can share memory with D

# Imperative Programs



```
import numpy as np  
a = np.ones(10)  
b = np.ones(10) * 2  
c = b * a  
print c  
d = c + 1
```

Easy to debug,  
easy to code

- **Advantages**

- Straightforward and flexible
- Use language natively (loops, control flow, debugging)

- **Disadvantages**

- Hard to optimize  
(Fixed with JIT compiler!)

# MXNet Gluon

Q: Why GLUON ?

A:



# Laptop Setup

- Get today's notebooks  
`git clone https://github.com/szha/KDD18-Gluon`
- Set up Conda environment
  - `conda env update --prune mxnet/kdd18_gluon_nlp`
  - `conda activate kdd18_gluon_nlp`



# Resources

- Deep Learning – the Straight Dope  
<https://gluon.mxnet.io/>
- A 60-minute Gluon Crash Course  
<https://gluon-crash-course.mxnet.io/>
- GluonCV  
<http://gluon-cv.mxnet.io/>
- GluonNLP  
<https://gluon-nlp.mxnet.io/>
- MXNet User Forum  
<http://discuss.mxnet.io/>
- MXNet Documentation  
<https://mxnet.apache.org/>



# Outline

<b>8:30-9:15</b>	Installation and Basics (NDArray, AutoGrad, Libraries)
<b>9:15-9:30</b>	Neural Networks 101 (MLP, ConvNet, LSTM, Loss, SGD) - Part I
<b>9:30-10:00</b>	Break
<b>10:00-10:30</b>	Neural Networks 101 (MLP, ConvNet, LSTM, Loss, SGD) - Part II
<b>10:30-11:00</b>	Computer Vision 101 (Gluon CV)
<b>11:00-11:30</b>	Parallel and distributed training
<b>11:30-12:00</b>	Data I/O in NLP (and iterators)
<b>12:00-13:30</b>	Break
<b>13:30-14:15</b>	Embeddings
<b>14:15-15:00</b>	Language models (LM)
<b>15:00-15:30</b>	Sequence Generation from LM
<b>15:30-16:00</b>	Break
<b>16:00-16:15</b>	Sentiment analysis
<b>16:15-17:00</b>	Transformer Models & machine translation
<b>17:00-17:30</b>	Questions



# NDArray and AutoGrad



# Matrices and Vectors

- A Simple ML algorithm

```
initialize function f(x,w)
```

```
for (x,y) in data:
```

```
    l = loss(y, f(x,w))
```

```
    g = gradient(l)
```

```
    w = w - eta * g
```

- Need vectors to deal with  $w$
- Need matrices to deal with chain rule for loss gradient
- Need high performance linear algebra



# NDArray (Linear Algebra on GPUs in Python)

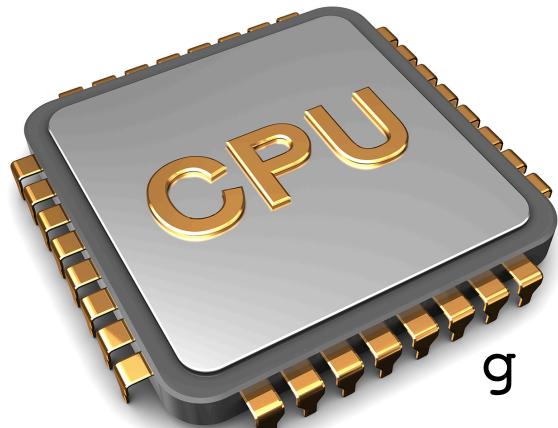
- NumPy
  - Limited to CPUs
  - No automatic differentiation
  - **Blocking (returns only once computation is performed)**
- NDArray
  - Multiple CPUs / GPUs via device context
  - Distributed systems in the cloud
  - Nonblocking Python frontend (no GIL)
  - Lazy evaluation

# NDArray Device Contexts

```
context = mx.cpu()
```

```
context = mx.gpu(0)
```

```
context = mx.gpu(1)
```



```
g = copyto(c)
```

```
g = c.as_in_context(mx.gpu(0))
```

# Automatic Differentiation

```
x = nd.array([[1, 2], [3, 4]])  
x.attach_grad()  
with ag.record():  
    y = x * 2  
    z = y * x  
z.backward()
```

$$\partial_x \sum_i z[i](x) \text{ not } \partial_x z[i](x)$$

**AutoGrad computes gradients  
without a declared compute graph**